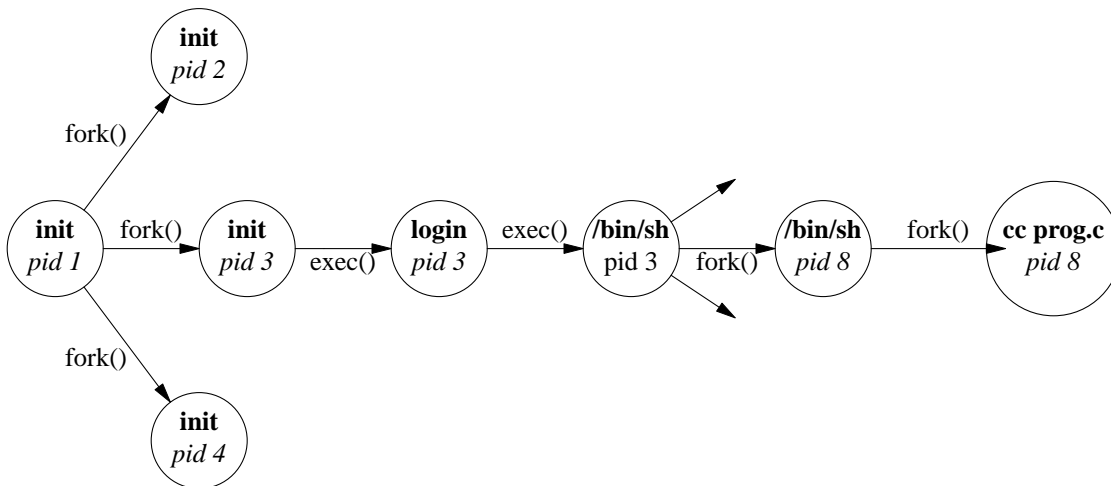


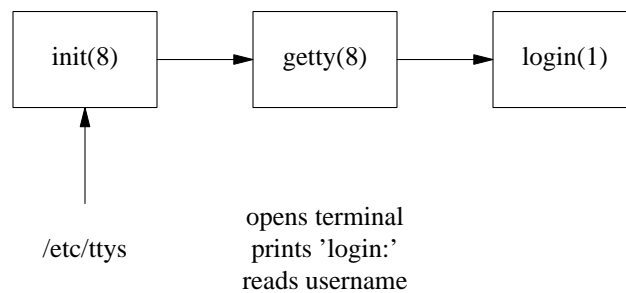
Unix Login Process

Ian Stanley

Init (PID 1) will fork for a number of times (at least once for each getty to be used). Each forked init then exec's to run login. Once a user enters his/her credentials the login execs to run the user's shell. The shell forks to run each command using a fork-exec combo.



The first few steps of the login process all run with superuser privileges.



login(1) runs the following:

- getpass(3), has, compare to getpwnam(3)
- if successful registers login in system databases

- reads/displays various files eg. motd
- a user may be part of a number of supplementary groups, so login initialises group membership: `initgroups(3)`, `setgid(2)`, initialize environment
- `chdir(2)` to new home directory
- `chown(2)` of the terminal device
- finally `setuid(2)` to user's uid and `exec(3)` a shell

The kernel explicitly creates `init(8)`.

`init(8)` creates `getty(8)` using `fork()` and `exec()`.

`getty(8)` execs (but does not fork) `login(1)`.

`login(1)` execs (but doesn't fork) `$SHELL`.

And the `$SHELL` forks and execs commands.

<code>init(8)</code>	pid 1, ppid 0, euid 0
<code>getty(8)</code>	pid #, ppid 1, euid 0
<code>login(1)</code>	pid #, ppid 1, euid 0
<code>\$SHELL(1)</code>	pid #, ppid 1, euid U
<code>ls(1)</code>	pid ##, ppid #, euid U

Where # is a unique PID no ≤ 1 and ## a unique PID \leq #